

Local Evidence Is Not Completion: Evidence-Condition Geometry for Bounded Completion-Audit Workflows

Anonymous Submission

Abstract

Completion-audit workflows often stop before the remaining workload is known. Their stop signals—no-new findings, agent agreement, stable summaries, or self-reported completion—can be locally meaningful while still being unsafe as scope-wide completion certificates. The missing question is: under what evidence condition was the stop signal produced, and does that condition match the claimed scope? We study this certificate mismatch in bounded source-route audit workflows. A source-route stratum records both where the workflow searched and which audit route it applied. We use the resulting exposure geometry to build an evidence-condition controller with four modules: source-route exposure estimation, an eligibility gate, residual-evidence repair over weak plausible gaps, and a final SAFE/CONTINUE/ABSTAIN decision rule. Across four bounded completion-audit task families, homogeneous route reuse creates false certification risk; source-route granularity detects this risk where source-only control does not; broad eligibility alone is insufficient on the `urllib3` boundary; and the full controller avoids unsafe stops without becoming a never-stop rule. The result is a bounded diagnostic and control principle: completion certificates should be judged by the evidence condition that produced them, not by stop signals alone.

1 Introduction

This paper evaluates bounded completion-audit workflows with declared source-route scopes. These are workflows whose intended evidence region is bounded, such as a repository snapshot or document set, and whose completion claim is stated over a declared set of sources and audit routes. Examples include searching a repository snapshot for all instances of a behavior, auditing a document set for policy exceptions, checking whether an implementation covers declared failure-mode routes, or deciding whether a multi-agent audit has exhausted useful evidence within its declared scope. Such workflows must eventually stop. The difficulty is that the absence of newly found evidence is not the same object as evidence that the intended bounded workload is complete.

The failure studied in this paper is *certificate mismatch*. A workflow may produce a plausible local stop signal: repeated scans return no new findings, agents agree with each other, subtasks appear closed, or a summary becomes stable. These signals are conditioned on the parts of the task that were

searched and on the routes by which they were searched. If the workflow repeatedly applies the same route to the same subset of sources, a no-new signal may certify only local exhaustion. It does not automatically support a global claim that the task is complete over its intended scope.

Consider a code-audit workflow. Auditing a file for timeout behavior does not certify that the same file has been audited for TLS behavior, exception behavior, retry behavior, or cleanup behavior. The source may be familiar, the agents may agree, and the workflow may have stopped finding timeout-related evidence; none of that establishes that other routes over the same source have been exhausted. The stop evidence is local in its evidence condition, while the completion claim is global in scope.

We formalize this observation with *evidence-condition geometry*. The basic unit is a source-route stratum, consisting of a bounded evidence region and an audit or search route. At runtime, the workflow accumulates exposure counts over these strata. The normalized exposure distribution describes the condition under which completion evidence was produced. When this distribution is localized, stop evidence is locally conditioned. It may be useful, but it is not a global completion certificate.

Based on this geometry, we define an evidence-condition controller. At a proposed stop, the controller computes the source-route exposure condition, checks whether the condition is broad enough to make the completion claim eligible for certification, and triggers repair over weak plausible strata when necessary. The controller is intentionally conservative: broad exposure gives completion eligibility, not automatic safety. A SAFE decision additionally requires that repair or audit reveal no residual evidence. If residual evidence appears, the controller outputs CONTINUE; if budget is exhausted without a defensible certificate, it outputs ABSTAIN.

Our experiments evaluate this controller on four bounded completion-audit task families: two generated scored tasks and two external real-repository audits with pattern-defined oracles. These tasks are not presented as a benchmark for the best item discovery policy. Instead, they test whether source-route exposure localization diagnoses unsafe stop claims and whether a controller can avoid accepting locally conditioned evidence as global completion proof.

This paper makes three contributions. First, it formulates false completion as a certificate mismatch between the declared scope of a stop claim and the source-route condition under which supporting evidence was produced. Second, it defines a certificate semantics for bounded audit completion:




Figure 1: Certificate mismatch from local stop evidence. A workflow may observe no-new findings, agreement, or a stable summary under evidence condition U . Promoting that stop proposal to a completion claim over Ω is unsafe when $U \neq \Omega$; the unobserved source-route strata are not certified.

a SAFE decision is admissible only when the source-route evidence condition is eligible and residual-evidence repair is nonproductive under runtime-visible signals. Third, it provides bounded audit evaluation, including generated audits and external repository audits, showing that homogeneous route reuse creates false certification risk while broader source-route evidence improves eligibility without being sufficient for safety.

We deliberately do not claim that residual-potential is an optimal repair method. It is one mechanism-aligned repair instance used to test whether weak source-route regions can expose residual evidence. The main claim is about the conditions under which completion evidence can support a stop decision.

2 Related Work

High-recall review, active search, and missing mass. Technology-assisted review and total-recall systems study how to retrieve nearly all relevant items under review budgets [Grossman and Cormack, 2011, Cormack and Grossman, 2014, 2015, 2016, Grossman et al., 2016, Roegiest et al., 2015], while active search and active learning choose actions or examples to expose positives efficiently [Lewis and Gale, 1994, Settles, 2009, Garnett et al., 2012, Attenberg et al., 2015]. Missing-mass and unseen-species estimators ask how much probability mass remains unseen [Good, 1953, Chao, 1984, Chao et al., 2014]. Our scored subclasses use recall only for post-hoc false-certification measurement. The controller does not observe oracle totals, recall, or hidden missing mass; it observes the source-route condition under which a stop claim was produced and whether repair reveals residual evidence.

Agent workflows and verification. Tool-using, browser-assisted, multi-agent, and memory-based language agents allocate work across tools, agents, searches, and intermediate states [Yao et al., 2023, Schick et al., 2023, Nakano et al., 2021, Qin et al., 2024, Wu et al., 2023, Park et al., 2023, Shinn et al., 2024, Madaan et al., 2023, Wang et al., 2024].

Agent benchmarks and factual verification benchmarks evaluate task success, tool use, factuality, or evidence support [Yao et al., 2022, Deng et al., 2024, Zhou et al., 2024, Liu et al., 2024, Jimenez et al., 2024, Hendrycks et al., 2021a, Liang et al., 2022, Srivastava et al., 2023, Lin et al., 2022, Thorne et al., 2018, Gao et al., 2023, Min et al., 2023, Schuster et al., 2021]. Our question is narrower and precedes final answer scoring: whether the evidence condition under which a workflow stopped is broad enough for the scope of its completion claim.

Agreement, audit agents, and oversight. Ensembles, self-consistency, verifier-guided systems, and audit agents can improve reliability [Dietterich, 2000, Cobbe et al., 2021, Wang et al., 2023]. Agreement is not useless, but it is unsafe as an unconditioned completion certificate: agents may agree because they repeatedly examined the same local region. Safety and oversight work studies broader failure modes and scalable supervision [Amodei et al., 2016, Christiano et al., 2017, Irving et al., 2018, Leike et al., 2018, Hendrycks et al., 2021b]. This paper targets a specific operational failure within that broader landscape: a workflow may present a global completion certificate whose supporting evidence was produced under a local source-route condition.

Coverage and stopping criteria. Software testing, audit coverage, and review stopping rules ask whether enough program structure, test paths, documents, or positives have been covered. Our use of source-route coverage is different: coverage is not the objective being maximized, but a diagnostic for certificate eligibility. We do not estimate how many items remain unseen, nor do we optimize active search for discovery yield. Instead, we ask whether the evidence condition that produced a stop certificate matches the scope that the certificate claims to cover.

3 Problem Setting

We study bounded completion-audit workflows with unknown residual workload and declared source-route scope. A workflow receives a task, allocates work across agents or tools, gathers evidence under a finite budget, and must decide whether the task is complete within that declared scope. Let

$$W = (T, S, R, A, B),$$

where T is the task, S is a set of sources, R is a set of routes, A is the set of agents or tools, and B is a finite budget. A source is a bounded evidence region, such as a file, module, document, source family, or repository component. A route is an audit or search lens applied to a source.

At runtime, the workflow produces evidence events

$$e_t = (a_t, x_t, r_t, u_t, o_t, c_t),$$

where a_t is the agent or tool, $x_t \in S$ is the source, $r_t \in R$ is the route, u_t is the action, o_t is the observation, and c_t is

the cost. Every evidence event is produced under a source-route condition. A source-route stratum is $s = (x, r)$, and the intended evidence-condition space is

$$\Omega = S \times R.$$

The source-route space does not require the workflow to know all true missing items in advance. It requires the intended audit scope to be defined: which sources and routes are relevant to the completion claim. This assumption is weaker than knowing the answer set and stronger than an unbounded open-world search. It matches many practical audits: a system may know which repositories, documents, modules, or evidence sources are in scope, and it may know which audit routes are relevant, while still not knowing how many true findings remain. The controller therefore evaluates the condition under which evidence was gathered, not the unknown total workload itself.

A stop claim is a workflow assertion

C_t : T is complete over the intended workload scope.

A false certification occurs when the workflow accepts SAFE for a stop claim that is not supported by the evaluation oracle:

$$\text{SAFE}(C_t) = \text{true}, \quad \text{complete}(T) = \text{false}.$$

In item-discovery subclasses, this is operationalized as accepting SAFE while bounded-oracle recall is below the completion threshold. In broader completion audits, it means accepting a global completion claim while relevant support, contradiction, or unresolved evidence remains outside the evidence condition. The oracle is used only for evaluation. It is not available to route assignment, repair scoring, eligibility checks, or stop decisions.

4 Evidence-Condition Geometry

Let $v_t(s)$ be the runtime-visible exposure count for source-route stratum s up to time t . Exposure may count visits, scans, route-specific audits, tool calls, or other logged search actions. Define

$$p_{\text{exp}}(t, s) = \frac{v_t(s)}{\sum_{s' \in \Omega} v_t(s')}.$$

This distribution is a point on the source-route coverage simplex. It describes where and how completion evidence was produced.

We use two runtime summaries:

$$\text{support}(t) = \frac{|\{s : v_t(s) > 0\}|}{|\Omega|},$$

$$G_{\text{exp}}(t) = \text{Gini}(\{v_t(s) : s \in \Omega\}).$$

These summaries are operational diagnostics, not universal laws. They summarize whether the evidence condition is localized or broad enough to make a completion certificate eligible for consideration.

Proposition 1 (Local no-new evidence is not a global certificate). *Let $U \subset \Omega$ be a strict subset of the intended source-route space. Suppose all evidence used to support a stop claim C_t is produced inside U , and suppose the workflow observes no-new evidence only for actions inside U . Then the no-new evidence can support local exhaustion over U , but it cannot by itself support global completion over Ω .*

Proof sketch. No-new evidence is conditioned on the source-route actions that were executed. If no events are produced in $\Omega \setminus U$, the workflow has not observed the absence of residual evidence in those unexposed strata. A world in which U is exhausted and $\Omega \setminus U$ contains residual evidence is observationally indistinguishable from a globally complete world under the workflow’s local evidence log. Therefore local no-new evidence rules out additional findings only under the searched condition U ; without additional assumptions about unsearched strata, it cannot rule out residual evidence in $\Omega \setminus U$.

Source-only geometry is too coarse: auditing a file for timeout behavior does not certify that TLS, retry, exception, or cleanup behavior has been audited in that file. Source-route-action geometry is more detailed, but it is sensitive to logging conventions and current experiments do not show a stable advantage over source-route. Source-route is the default because it captures both where and how evidence was produced while remaining interpretable and runtime-computable. The geometry is thus a certificate diagnostic rather than a semantic theory of all possible evidence. It asks whether the current log covers the intended source-route space broadly enough for the proposed stop claim; it does not assert that all routes are equally important or that support/Gini thresholds transfer unchanged across domains.

5 Method

The method has four modules: source-route exposure estimation, eligibility gating, residual-evidence repair, and the final SAFE/CONTINUE/ABSTAIN decision rule. The point is not to prove universal completion; it is to make the runtime evidence condition explicit before a stop claim is accepted.

5.1 Source-Route Exposure Estimation

Input: a runtime evidence log E_t , intended source set S , intended route set R , and their product space $\Omega = S \times R$. Output: exposure counts $v_t(s)$, normalized exposure $p_{\text{exp}}(t, s)$, source-route support $\text{support}(t)$, and exposure concentration $G_{\text{exp}}(t)$. Role: summarize where and how evidence was gathered, so localized stop evidence can be distinguished from broad evidence. Algorithm 1 shows how the four modules are composed at a proposed stop.

The source-route exposure summaries are operational diagnostics, not universal laws. They summarize whether the evidence condition is localized or broad enough to make a completion certificate eligible for consideration.




Figure 2: Evidence-condition controller. The controller maps a runtime evidence log to an exposure distribution over the intended source-route space. The support/Gini check is an eligibility test, not a completion proof. If the condition is localized or weak plausible gaps remain, the controller repairs or continues. It returns SAFE only when the evidence condition is broad and repair or audit reveals no residual evidence.

Algorithm 1 Evidence-Condition Controller

Require: log E_t , space Ω , claim C_t , budget B , thresholds τ_s, τ_g , repair rule $\mathcal{R}$

Ensure: decision in {SAFE, CONTINUE, ABSTAIN}

- 1: Compute $v_t(s)$, $p_{\text{exp}}(t, s)$, $\text{support}(t)$, and $G_{\text{exp}}(t)$
 - 2: $\text{eligible} \leftarrow [\text{support}(t) \geq \tau_s \wedge G_{\text{exp}}(t) \leq \tau_g]$
 - 3: **if** not eligible and $B = 0$ **then**
 - 4: **return** ABSTAIN
 - 5: **end if**
 - 6: **if** not eligible or weak plausible gaps remain **then**
 - 7: Audit $Q = \mathcal{R}(E_t, \Omega, C_t)$ within B
 - 8: **if** residual evidence appears **then**
 - 9: **return** CONTINUE
 - 10: **end if**
 - 11: Recompute eligible
 - 12: **end if**
 - 13: **if** eligible and no residual evidence appeared **then**
 - 14: **return** SAFE
 - 15: **end if**
 - 16: **return** ABSTAIN
-

5.2 Eligibility Gate

Input: $v_t(s)$, $p_{\text{exp}}(t, s)$, $\text{support}(t)$, $G_{\text{exp}}(t)$, thresholds τ_s and τ_g . Output: eligible / ineligible. Role: decide whether the current evidence condition is broad enough for a stop claim to be considered.

The eligibility gate is not a completion proof. It is only a test of whether the evidence condition is broad enough for a global stop claim to be considered. In the current implementation, support and Gini thresholds instantiate this operational gate; they are not claimed as universal laws.

5.3 Residual-Evidence Repair

Input: current evidence log E_t , evidence condition Ω , claim C_t , budget B , and repair rule $\mathcal{R}$. Output: repaired log or continued

audit targets. Role: probe weak plausible strata for residual evidence when eligibility alone is insufficient. Weak plausible gaps are strata that are both under-exposed in the current source-route condition and still plausible under non-oracle runtime signals such as source text, route names, route match counts, source size, or ledger-visible agent notes.

The repair instance used in the experiments is

$$\text{priority}(s) = \text{under_exposure}(s) \cdot \text{runtime_potential}(s).$$

The under-exposure term gives priority to weak parts of the current completion certificate. The runtime-potential term uses visible information such as source text, route names, source size, route match counts, or non-oracle ledger signals. It is not allowed to use oracle labels, oracle totals, oracle missing mass, undiscovered true item counts, post-hoc recall, or scorer-visible target distributions. Residual-potential is a mechanism-aligned repair instance, not an optimal active-search algorithm.

5.4 SAFE/CONTINUE/ABSTAIN Decision Rule

Input: eligibility result and residual-evidence repair outcome. Output: SAFE, CONTINUE, or ABSTAIN. Role: accept a stop only when the evidence condition is broad and repair or audit reveals no residual evidence.

The final rule is intentionally conservative. Broad exposure gives completion eligibility, not automatic safety. A SAFE decision additionally requires that repair or audit reveal no residual evidence. If residual evidence appears, the controller outputs CONTINUE; if budget is exhausted without a defensible certificate, it outputs ABSTAIN.

Proposition 2 (Controller certificate semantics). *For a declared source-route scope Ω , thresholds τ_s, τ_g , and a fixed repair rule $\mathcal{R}$ that uses only runtime-visible signals, Algorithm 1 returns SAFE only if (i) the exposure condition satisfies the eligibility gate and (ii) repair or audit reveals no residual evidence within the allocated repair budget. If residual evidence is found, the returned decision is CONTINUE; if eligibility*

cannot be established under budget, the returned decision is ABSTAIN.

Proof sketch. The algorithm computes eligibility before any SAFE branch is reachable. Ineligible states either enter repair or return ABSTAIN when no budget remains. During repair, any residual evidence triggers the CONTINUE branch. The only SAFE branch is guarded by both eligibility and the absence of residual evidence. The statement is therefore a soundness property of the controller’s declared certificate semantics, not a guarantee that the world contains no unobserved evidence.

6 Experiments

The experiments test whether a proposed stop can be accepted as a bounded completion certificate under a declared source-route scope. They are not a benchmark for the best item-discovery policy. Each run follows the same closed protocol. First, a workflow trajectory produces a proposed stop state using only runtime-visible evidence. Second, the compared decision rule observes the exposure log, applies any allowed repair under the fixed budget, and returns SAFE, CONTINUE, or ABSTAIN. Third, the oracle is revealed only for post-hoc scoring. The controller is therefore evaluated on whether it accepts unsafe completion claims, not on whether it has oracle access to the remaining workload.

We organize the evaluation around five method-aligned questions. Q1 asks whether naive stop signals create false certification. Q2 asks whether source-route geometry detects risks that source-only exposure misses. Q3 asks whether broad eligibility is sufficient for safety. Q4 asks whether the full controller rejects unsafe stops without degenerating into a never-stop rule. Q5 compares residual-potential repair with random and high-potential repair under the same decision rule.

Tasks. `policy-docset` is a generated policy document set with four sources covering access control, data handling, release process, and audit exceptions. Its routes are obligation, exception, deadline, and prohibition; the hidden oracle contains 24 policy clauses. `code-repo` is a generated bounded code repository with authentication, payment, storage, and API-client sources. Its routes are compatibility, security, and resilience; the hidden oracle contains 20 code-risk items.

`requests` is an external real-repository audit over six files from a local installed snapshot of the Python package: adapters, API, authentication, models, sessions, and utilities. Its routes are TLS, timeout, exception, and compatibility. `urllib3` is an external real-repository audit over connection, connection-pool, pool-manager, response, retry, and timeout sources. Its routes are timeout, retry, TLS, exception, and cleanup, giving 30 intended source-route strata. External oracles are pattern-defined from frozen snapshots: stronger than generated toys, weaker than human-annotated completion-audit benchmarks.

Workflow conditions. Homogeneous route reuse assigns multiple agents to the same route across sources. The reused

routes are obligation for `policy-docset`, compatibility for `code-repo`, and timeout for the two external repository audits. Route-partitioned audit assigns agents to different routes over the same source set. `urllib3` additionally includes an `extended_audit` condition covering all five routes.

Compared policies and ablations. We compare decision variants and repair-target variants on fixed bounded-audit states. Decision variants are Naive stop, Source-only, Source-route eligibility-only, and the full controller. Repair-target variants keep the full decision rule fixed and choose strata by Random, High-potential, or Residual-potential. All policies use the same fixed trajectories and proposed stop states; differences come from the decision rule or repair target selection, not from different discovery trajectories. Oracle labels are used only for post-hoc scoring.

Seeds, budgets, thresholds, and cost. Generated tasks use repair budget 4; `requests` and `urllib3` use 200 seeded validations with budgets 4 and 5, respectively. The controller uses $\tau_s = 0.75$, $\tau_g = 0.70$ as pre-specified operational test points, not oracle-tuned hyperparameters. A recall threshold of 0.90 is used only for bounded-oracle false-certification evaluation and is invisible to the controller. Repair cost counts only additional scanned source lines plus extraction events after a stop state is fixed. Seeds, route assignments, repair orders, and challengers are fixed before oracle scoring.

Reproducibility and leakage control. We provide both a lightweight three-seed reproduction configuration and the full seeded validation configuration. The configs specify task snapshots, route sets, thresholds, budgets, seeds, repair rules, oracle-scoring paths, output CSVs, and figure paths. Runtime decisions may use only the evidence log, source text, lexical route matches, route identifiers, source size, and non-oracle ledger notes. They may not use oracle totals, undiscovered item counts, post-hoc recall, or scorer-visible target distributions. The released scripts regenerate Table 1, Figure 3, Figure 4, and the controller-count tables from the same fixed configs.

The four task families supply bounded completion-audit cases with withheld or post-hoc oracles. Simulation sweeps are screening checks, not new task claims. Leakage controls are in the supplement; runtime decisions use visible logs and source text, not oracle labels or undiscovered true-item counts.

7 Results

Naive stop signals cause false certification. Figure 3 begins with the failure mode. Across all four tasks, homogeneous route reuse produces localized evidence and would falsely certify completion if accepted as SAFE. Controlled sweeps in the supplement show the same qualitative pattern under localized exposure. This motivates source-route granularity rather than accepting unconditioned stop signals.

Table 1: Experimental coverage. SR strata denote source-route strata; oracles are used only for post-hoc scoring.

Study	Sources	Routes	SR strata	Oracle items	Runs/seeds	Main measurement
policy-docset	4	4	16	24	200/challenger	FCR, repair gain
code-repo	4	3	12	20	200/challenger	FCR, repair gain
requests	6	4	24	298	200/challenger	controller validation
urllib3	6	5	30	699	200/challenger	boundary validation
Controlled geometry	–	variable	variable	synthetic	2160	metric screening
Exposure localization	–	variable	18/30/42	synthetic	2700	localization check




Figure 3: Main diagnostic and ablation. Left: broader audits increase source-route support. Middle: source-only support is full in homogeneous runs, but source-route support remains localized. Right: the eligible `urllib3` route-partitioned state still has recall 0.835, below the post-hoc 0.90 threshold, while the extended audit reaches recall 1.000.

Table 2: Boundary case showing why eligibility is not proof. The state passes the source-route eligibility gate, but post-hoc recall remains below the completion threshold; repair-aware control therefore returns CONTINUE rather than SAFE.

Task	State	Support	Gini	Recall	Residual+	Elig.-only	Full
urllib3	route_partitioned	0.800	0.647	0.835	yes	SAFE	CONTINUE

Source-route geometry is necessary. The middle panel of Figure 3 is the source-only vs. source-route ablation. Homogeneous runs touch all files, so source-only support looks complete; the same states have localized source-route support and sub-threshold recall. Source-only exposure cannot detect that timeout evidence does not certify TLS, retry, exception, or cleanup routes in the same file.

Eligibility alone is not proof. Route-partitioned evidence improves eligibility, but it is not proof. In `urllib3`, route-partitioned exposure covers 24/30 strata and passes the gate, yet recall remains below threshold and repair finds residual evidence; the controller returns CONTINUE, not SAFE. The extended audit reaches support and recall 1.000.

Full controller avoids unsafe stops without never-stopping. A naive controller falsely certifies homogeneous stops; the

Table 3: Seeded controller decision counts for Figure 4. All decision rules are evaluated on the same seeded stop-boundary states; oracle labels are used only after decisions are fixed. S/C/A reports SAFE/CONTINUE/ABSTAIN counts.

Decision rule	Unsafe	Complete	Unsafe S/C/A	Complete S/C/A	FCR	Safe cov.
Naive stop	400	1200	400/0/0	1200/0/0	1.000	1.000
Source-only	400	1200	400/0/0	1200/0/0	1.000	1.000
Eligibility-only	400	1200	0/0/400	1200/0/0	0.000	1.000
Full	400	1200	0/400/0	1200/0/0	0.000	1.000

full controller rejects them. It is not a never-stop rule: it returns SAFE for broad complete states in the generated tasks, `requests`, and the extended `urllib3` audit. In seeded external validations, repair states are oracle-unsafe and return CONTINUE or ABSTAIN, while oracle-safe complete-state perturbations keep safe coverage 1. Figure 4 therefore reports seeded unsafe-state FCR rather than relying only on the five fixed observed stop states; those fixed states define the boundary cases and are counted in the supplement.

Repair policy comparison. Figure 4 separates repair target variants from decision variants. Residual-potential exposes residual evidence and has higher mean gain than random repair, but it is not proven optimal; high-potential is competitive in cost-normalized terms. Repair is therefore a bounded probe




Figure 4: Controller variant comparison. Left/middle: decision variants are scored on seeded unsafe repair states and seeded complete-state perturbations derived from the same fixed stop boundaries. The decision distribution panels show SAFE, CONTINUE, and ABSTAIN rates on unsafe and complete states; Table 2 isolates the eligible-but-residual-positive boundary case. Right: repair target variants keep the full decision rule fixed and compare mean repair gain vs. mean post-stop repair cost over seeded unsafe repair states.

for remaining evidence, not a claim of optimal active search. The supplement includes per-task decision breakdowns and lightweight external baselines, including source-only stopping, random route expansion, high-potential expansion, and a scalar singleton/Chao negative control.

Sensitivity and audit checks. Sensitivity sweeps are in the supplementary material. Varying τ_s, τ_g and repair budget changes the SAFE/CONTINUE/ABSTAIN mixture and repair cost, but does not create false certifications in the external controller validations. Oracle construction, negative controls, leakage checks, seeded details, and regex examples are also supplementary.

Together, the results close the intended bounded-audit loop: naive and source-only stopping certify localized evidence, source-route geometry exposes the mismatch, eligibility is only a precondition, and repair distinguishes unsafe residual states from broad complete states.

8 Discussion

The experiments support a bounded certificate claim rather than an open-world guarantee. A SAFE decision means that the declared source-route condition is eligible and that budgeted repair did not reveal residual evidence; it does not mean that the world contains no further facts or that future searches would fail.

What the validation establishes. The ablations close the intended loop: homogeneous route reuse induces unsafe local stop evidence; source-route exposure detects the mismatch that source-only coverage misses; the `urllib3` boundary shows that eligibility is only a precondition; and repair separates broad-but-productive states from clean complete states.

Reproducibility and auditability. The implementation fixes snapshots, routes, seeds, thresholds, budgets, and oracle-

scoring paths in config files. A three-seed config checks regeneration; the full seeded config produces the reported validation. Runtime-visible fields are separated from oracle-only fields, and stop states are fixed before scoring, so the controller cannot use the hidden workload it is supposed to audit.

9 Limitations

The external repository oracles are pattern-defined and therefore stronger than generated toy labels but weaker than human-annotated completion-audit benchmarks. They are nevertheless appropriate for this mechanism validation because the evaluation asks whether the controller accepts unsafe completion certificates, not whether it performs semantic vulnerability discovery. Source-route strata are task-designed, and other domains may need different route inventories or thresholds. Residual-potential repair is a mechanism-aligned probe, not an optimal active-search policy. The workflow pilot is an interface validation rather than a scored claim-verification benchmark; a human-annotated completion-claim audit remains the key external-validity test.

10 Conclusion

Local evidence does not automatically support a scope-wide completion claim. Evidence-condition geometry makes this conditioning explicit: the controller rejects unsafe local stops, repairs weak plausible strata, and abstains when no defensible certificate can be formed. Completion certificates should be judged by evidence condition, not stop signals alone.

References

- Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete problems in ai safety. *arXiv preprint arXiv:1606.06565*, 2016.
- Josh Attenberg, Panos Ipeirotis, and Foster Provost. Beat the machine: Challenging humans to find a predictive model’s unknown unknowns. *Journal of Data and Information Quality*, 6(1), 2015.
- Anne Chao. Nonparametric estimation of the number of classes in a population. *Scandinavian Journal of Statistics*, 11(4):265–270, 1984.
- Anne Chao, Nicholas J. Gotelli, T. C. Hsieh, Elizabeth L. Sander, K. H. Ma, Robert K. Colwell, and Aaron M. Ellison. Rarefaction and extrapolation with hill numbers: A framework for sampling and estimation in species diversity studies. *Ecological Monographs*, 84(1):45–67, 2014.
- Paul F. Christiano, Jan Leike, Tom B. Brown, Miljan Martić, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems*, 2017.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Gordon V. Cormack and Maura R. Grossman. Evaluation of machine-learning protocols for technology-assisted review in electronic discovery. In *Proceedings of the 37th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 153–162, 2014.
- Gordon V. Cormack and Maura R. Grossman. Multi-faceted recall of continuous active learning for technology-assisted review. In *Proceedings of the 38th International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2015.
- Gordon V. Cormack and Maura R. Grossman. Engineering quality and reliability in technology-assisted review. *Proceedings of the 39th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 75–84, 2016.
- Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen, Sam Stevens, Boshi Wang, Huan Sun, and Yu Su. Mind2web: Towards a generalist agent for the web. In *Advances in Neural Information Processing Systems*, 2024.
- Thomas G. Dietterich. Ensemble methods in machine learning. *Multiple Classifier Systems*, pages 1–15, 2000.
- Luyu Gao, Zhuyun Dai, Panupong Pasupat, Anthony Chen, Arun Tejasvi Chaganty, Yicheng Fan, Vincent Y. Zhao, Ni Lao, Hongrae Lee, Da-Cheng Juan, and Kelvin Guu. Rarr: Researching and revising what language models say, using language models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, 2023.
- Roman Garnett, Yamuna Krishnamurthy, Xuehan Xiong, Jeff Schneider, and Richard Mann. Bayesian optimal active search and surveying. In *Proceedings of the 29th International Conference on Machine Learning*, pages 843–850, 2012.
- I. J. Good. The population frequencies of species and the estimation of population parameters. *Biometrika*, 40(3–4): 237–264, 1953.
- Maura R. Grossman and Gordon V. Cormack. Technology-assisted review in e-discovery can be more effective and more efficient than exhaustive manual review. *Richmond Journal of Law and Technology*, 17(3), 2011.
- Maura R. Grossman, Gordon V. Cormack, and Adam Roegiest. Trec 2015 total recall track overview. In *Proceedings of The Twenty-Fourth Text REtrieval Conference*, 2016.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *International Conference on Learning Representations*, 2021a.
- Dan Hendrycks, Nicholas Carlini, John Schulman, and Jacob Steinhardt. Unsolved problems in ml safety. *arXiv preprint arXiv:2109.13916*, 2021b.
- Geoffrey Irving, Paul Christiano, and Dario Amodei. Ai safety via debate. *arXiv preprint arXiv:1805.00899*, 2018.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. Swbench: Can language models resolve real-world github issues? In *International Conference on Learning Representations*, 2024.
- Jan Leike, David Krueger, Tom Everitt, Miljan Martić, Vishal Maini, and Shane Legg. Scalable agent alignment via reward modeling: A research direction. *arXiv preprint arXiv:1811.07871*, 2018.
- David D. Lewis and William A. Gale. A sequential algorithm for training text classifiers. In *Proceedings of the 17th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 3–12, 1994.
- Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, Benjamin Newman, Binhang Yuan, Bobby Yan, Ce Zhang, Connor Cosgrove, Christopher D. Manning, Christopher Ré, Diana Acosta-Navas, Drew A. Hudson, Eric Zelikman, Esin Durmus, Faisal Ladhak, Frieda Rong, Hongyu Ren, Huaxiu Yao, Jue Wang, Keshav Santhanam, Laurel Orr, Lucia Zheng,

- Mert Yuksekgonul, Mirac Suzgun, Nathan Kim, Neel Guha, Niladri Chatterji, Omar Khattab, Peter Henderson, Qian Huang, Ryan Chi, Sang Michael Xie, Shibani Santurkar, Surya Ganguli, Tatsunori Hashimoto, Thomas Icard, Tianyi Zhang, Vishrav Chaudhary, William Wang, Xuechen Li, Yifan Mai, Yuhui Zhang, and Yuta Koreeda. Holistic evaluation of language models. *arXiv preprint arXiv:2211.09110*, 2022.
- Stephanie Lin, Jacob Hilton, and Owain Evans. Truthfulqa: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*, 2022.
- Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, Minlie Huang, and Jie Tang. Agentbench: Evaluating llms as agents. In *International Conference on Learning Representations*, 2024.
- Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Sean Welleck, Bodhisattwa Prasad Majumder, Shashank Gupta, Amir Yazdanbakhsh, and Peter Clark. Self-refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems*, 2023.
- Sewon Min, Kalpesh Krishna, Xinxu Lyu, Mike Lewis, Wentau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. Factscore: Fine-grained atomic evaluation of factual precision in long form text generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 2023.
- Reiichiro Nakano, Jacob Hilton, Suchir Balaji, Jeff Wu, Long Ouyang, Christina Kim, Christopher Hesse, Shantanu Jain, Vineet Kosaraju, William Saunders, Xu Jiang, Karl Cobbe, Tyna Eloundou, Gretchen Krueger, Kevin Button, Matthew Knight, Benjamin Chess, and John Schulman. Webgpt: Browser-assisted question-answering with human feedback. *arXiv preprint arXiv:2112.09332*, 2021.
- Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, 2023.
- Yujia Qin, Shengding Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, Sihan Zhao, Yulong Hong, Ruobing Tian, Ruixue Xie, Jie Zhou, Mark Gerstein, Dahai Li, Zhiyuan Liu, and Maosong Sun. Toolllm: Facilitating large language models to master 16000+ real-world apis. In *International Conference on Learning Representations*, 2024.
- Adam Roegiest, Gordon V. Cormack, and Maura R. Grossman. Trec 2015 total recall track guidelines and assessment process. In *Proceedings of The Twenty-Fourth Text REtrieval Conference*, 2015.
- Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems*, 2023.
- Tal Schuster, Adam Fisch, and Regina Barzilay. Get your vitamin c! robust fact verification with contrastive evidence. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics*, 2021.
- Burr Settles. Active learning literature survey. Technical Report 1648, University of Wisconsin–Madison, 2009.
- Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems*, 2024.
- Aarohi Srivastava et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. In *Transactions on Machine Learning Research*, 2023.
- James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. Fever: A large-scale dataset for fact extraction and verification. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics*, 2018.
- Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. In *Transactions on Machine Learning Research*, 2024.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. In *International Conference on Learning Representations*, 2023.
- Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, Ahmed H. Awadallah, Ryen W. White, Doug Burger, and Chi Wang. Autogen: Enabling next-gen llm applications via multi-agent conversation. *arXiv preprint arXiv:2308.08155*, 2023.
- Shunyu Yao, Rohan Rao, Matthew Hausknecht, and Karthik Narasimhan. Alfworld: Aligning text and embodied environments for interactive learning. In *International Conference on Learning Representations*, 2022.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing

reasoning and acting in language models. In *International Conference on Learning Representations*, 2023.

Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. Webarena: A realistic web environment for building autonomous agents. In *International Conference on Learning Representations*, 2024.